

RISHABH SABHARWAL

+44 7486 045605 | r.sabharwal@sms.ed.ac.uk

<https://sabharwalrishabh.github.io/> | <https://www.linkedin.com/in/sabharwalrishabh/>

EDUCATION

MSc in Artificial Intelligence, University of Edinburgh, UK (2026) | **Grade: Distinction***

B.Tech in ECE (AI & ML), NSUT, India (2024) | **GPA: 3.68/4** (CVSPK Fellow, **Top 5%**)

SELECTED PUBLICATIONS

Rishabh Sabharwal, H. Wang, A. Storkey, J. Pan, “Multi-turn Evaluation of Deep Research Agents Under Process-level Feedback,” in SCALE Workshop, ICML 2026. (**Oral - Top 6%**)

A.Mazumder, Rishabh Sabharwal, ..., P. Rathore, “On Convergence of Adam with Data Dependent Stepsize” in **IEEE Transactions on AI** (Q1 Journal, Impact Factor: 4.8). ([Link](#))

K. Annuranjana, N. Shrivastava*, A. Johar*, Rishabh Sabharwal*, ..., “Does Reasoning Kill the Joke? Long-Context Humor Understanding in Hindi” in C3NLP Workshop, ACL 2026.

SELECTED PROJECTS

Multi-Turn Evaluation of Deep Research Agents (DRAs)

- Designed a novel multi-turn evaluation framework to assess agents’ ability to iteratively refine reports.
- Evaluations revealed a **major flaw in current DRAs**: models either regress severely or the models that tend to preserve content implicitly lead to 3-4x higher costs making multi-turn deep research infeasible.

LLM Post-Training on Small-Scale Models ([Code](#))

- Fine-tuned Qwen2.5-0.5B on GSM8K via **SFT** and **RLHF**; mitigated reward over-optimization using process reward models assigning dense rewards to CoTs, **improving accuracy from 31% to 37%** on test set
- Designed a **latent-space knowledge distillation** framework transferring K latent reasoning tokens from Qwen2.5-7B directly into the student’s representation space, achieving **41% accuracy** on the test set.

Multi-Step Agentic RAG System for Multi-Hop QA ([Code](#))

- Built a RAG pipeline for multi-hop QA over 270K facts using FAISS, HuggingFace embedding models, and reranking; served via FastAPI with vLLM-served Qwen3-8B and GPT-4.1-mini for inference.
- Designed a decoupled **agentic RAG system** with multi-step retrieval via LLM tool calling, improving EM from 0.56 (single-step RAG) to 0.67. Conducted systematic evaluation and error analysis using LLM-as-judge.

EXPERIENCE

ML Engineer

Jan 2024 - Jul 2025

Indian Institute of Science

Bengaluru, India

- Developed ReGA, a parameter-free spectral GNN module for global context filtering; achieved SOTA on 20 real-world datasets, reducing parameters by **55x** and GFLOPs by **100x**.
- Designed a plug-and-play graph-diffusion-inspired attention module for ConvNets. Our module **adds just 20 parameters to ResNet-50 yet outperforms ResNet-101** (~20M more parameters); generalizes across ConvNeXt-V2 and MogaNet backbones.

RELEVANT COURSEWORK

- Advanced NLP** (Transformers, PEFT, Knowledge Distillation, RAGs, Information Retrieval, Multimodal LLMs, LLM Agents), **ML Systems** (GPU Programming, LLM Inference Optimisation), **Probabilistic Modelling & Reasoning** (Bayesian Inference, Statistical ML, Optimisation Theory)

SKILLS

- Programming & Frameworks:** Python, Pandas, NumPy, PyTorch, Triton, FastAPI
- ML/LLM Tooling:** Hugging Face, PEFT (LoRA, QLoRA), vLLM, Weights & Biases, GCP (Vertex AI)
- Agentic AI:** LangChain, Multi-Agent Systems, Tool Calling
- DevOps/Systems:** Git, Linux, CI/CD, Distributed Training (PyTorch DDP)